

4. *ML project design with supervised learning (K-Nearest Neighbors)*

M.A.Z. Chowdhury and M.A. Oehlschlaeger

Department of Mechanical, Aerospace and Nuclear Engineering
Rensselaer Polytechnic Institute
Troy, New York

chowdm@rpi.edu
oehlsm@rpi.edu

"Until you spread your wings, you'll have no idea how far you can fly."
- Napoleon

MANE 4962 and 6962

End-to-end ML Project Design

- 1 Frame the problem and look at the big picture.
- 2 Get the data.
- 3 Discover and visualize the data to gain insights.
- 4 Prepare the data for Machine Learning algorithms.
- 5 Select a model and train it.
- 6 Fine-tune your model.
- 7 Present your solution.
- 8 Launch, monitor, and maintain your system.

There are other ways to design ML projects but this is a good starting point. Every step is described as a checklist in Appendix B of reference Aurélien Géron's book.

Working with Real Data

1 Popular open data repositories.

- UC Irvine Machine Learning Repository - Kaggle datasets
- Amazon's AWS datasets

2 Meta portals (they list open data repositories) — Data Portals

- OpenDataMonitor
- Quandl

3 Other pages listing many popular open data repositories

- Wikipedia's list of Machine Learning datasets
- Quora.com
- The datasets subreddit
- Github, Zenodo
- <https://www.uco.es/kdis/mlresources/>

Managing your project

Use version control to manage your code. I recommend GitHub. Following commands would come in handy.

```
git status
git add .
git commit -m "commit message"
git push
git pull
git clone
git init
git config
```

Data Wrangling with Pandas

Use the Pandas dataframe to work with your data.

```
import pandas as pd
df = pd.DataFrame()
df['column_name'] = a #a can be list/numpy array
df.head()
```

Help: <https://pandas.pydata.org/docs/reference/index.html>

Useful functions

head(), read_csv(), head(),
describe(), memory_usage(),
loc(), astype(), merge(), sort_values()

Sample project


- 1 We want to make a autonomous drone which is going to identify iris species.
- 2 It is going to measure the length and width of the sepals and petals.
- 3 Use the Iris dataset to design a simple machine learning model.


Similarity in nature

Birds of a feather flock together




A man or a horse or a horseman or a 'man-horse'?











k-Nearest Neighbors (kNN)


- 🔊 Classify or regress data points based on similarity to nearby data points.
- 🔊 Distance, $d(\underline{x}, \underline{x}') = ||\underline{x} - \underline{x}'||$
- 🔊 We can use L1 and L2 norms as distance metric to find nearby data points.
- 🔊 These are commonly known as taxicab or manhattan distance and the euclidean distance.
- 🔊 Scikit has kNNClassifier and kNNregressor classes for implementation.
- 🔊 For classification, the class associated with the nearest neighbors or the mean class is returned.
- 🔊 For regression, local interpolation of the targets associated with the nearest neighbors in the training set is returned.

k-NN algorithm (in simple words)

- 1 Find all the nearest neighbors
- 2 Show their average class as the class label of the unknown vector.

A visual example

- Say, we make three new measurements (×) and compare them to Fisher's Iris data.
- Predict the class labels (species types) for these measurements.



A simple implementation of K-NN


```

1  class My_KNNClassifier:
2      def __init__(self, k=3):
3          self.k = k
4      def fit(self, X_train, y_train):
5          self.X_train = X_train
6          self.y_train = y_train
7      def predict(self, X_test):
8          predictions = []
9          for i in range(X_test.shape[0]):
10             predictions.append(self._knn_classifier(X_test[i]))
11          return predictions
12      def _knn_classifier(self, X_test):
13          distances, targets = [], []
14          for i in range(self.X_train.shape[0]):
15              distance = np.linalg.norm(self.X_train[i]-X_test)
16              distances.append([distance, i])
17          distances = sorted(distances)
18          for i in range(self.k):
19              index = distances[i][1]
20              targets.append(self.y_train[index])
21          return max(targets, key=targets.count)
22  model = My_KNNClassifier()
23  model.fit(X_train, y_train)
24  preds = model.predict(X_test)
25  print(accuracy_score(y_test, preds))

```


All the calculations are inside predict method. See notebook for comparison and tuning.

Voronoi Diagram



Voronoi Diagram


- ➦ A geometric representation of a dataset that partitions a space into regions based on the distance to a set of points.
- ➦ The points in the dataset are called the **generators** or **seeds** of the Voronoi diagram, and each region is defined as the set of points that are closest to a particular generator.
- ➦ Each region is represented by a polygon, which is the set of points that are closer to the generator of that region than to any other generator.
- ➦ **Voronoi edges:** Edges of the polygons are the lines that divide the space into the different regions.
- ➦ **Voronoi vertices:** The points where the Voronoi edges meet. These points are equidistant to two or more generators.



For two iris features.

Voronoi Diagram

- Interpreting a Voronoi diagram can provide insight into the structure of a dataset by highlighting patterns and relationships between the generators.
- If the generators are evenly spaced, the resulting Voronoi diagram will have regular, symmetric regions.
- If the generators are clustered together, the resulting diagram will have irregular regions with many small polygons.
- It is useful to visualize the k-NN decision boundary that separate the classes, where the decision boundary is the set of points that are equidistant to two or more class generators.
- This helps understand the shape and complexity of the decision boundary, and to identify any potential issues with overfitting or underfitting.







For two iris features.

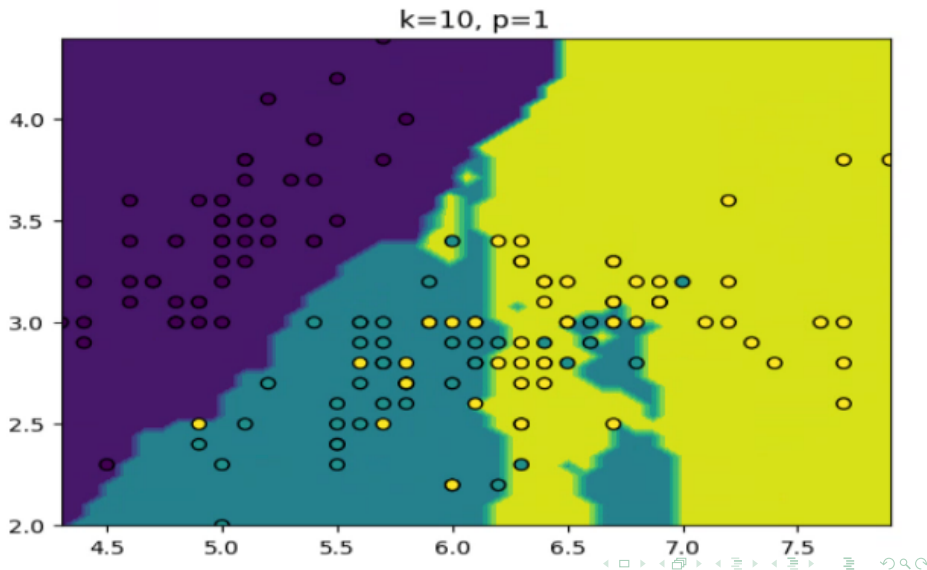
Choosing k

- 👉 $k=3$ [Good starting point for many problems]
- 👉 $k = \sqrt{N}$, N is the size of the train set.
- 👉 Cross-validation
 - small k
 - 👉 good for capturing fine-grained patterns
 - 👉 but it may overfit since it will be sensitive to random peculiarities present in the training data
 - large k
 - 👉 makes stable predictions by averaging over lots of examples
 - 👉 but it may underfit by failing to capture important regularities

Varying k



Varying k



Remarks

● Advantages:

- ☞ Fast learning, no explicit training
- ☞ Simple, easy to explain the method and results to users and customers.
- ☞ Can control the complexity by varying k
- ☞ Number of computations at training time is zero!
- ☞ Can easily do multi-class.
- ☞ Can easily adapt to regression.

● Disadvantages:

- ☞ Memory intensive since it needs to store the entire dataset in memory!
- ☞ Predictions can take a long time
- ☞ No model to shed light on process that generated data
- ☞ Thousands of work hours have gone into algorithms and data structures for efficient nearest neighbors with high dimensions and/or large datasets.
- ☞ **Suffers from the Curse of Dimensionality**

Applicable to high dimensional data




Image \rightarrow Matrix

0	10	13	0
0	0	22	0
0	31	0	0
0	11	15	7

0
10
13
0
0
0
22
0
0
31
0
0
0
0
11
15
7

\rightarrow Vector

Advantages of using feature vector or input vector

- ML algorithms handles lots of different types of data: signals, images, text, video etc.
- Represent the input as an input vector $\in \mathbb{R}^{dim}$
- **Representation** = mapping to another space that's easy to manipulate.
- Vectors provide a wonderful mechanism to work with the input features of the data because we can use linear algebra!