

16. Support Vector Machine

M.A.Z. Chowdhury and M.A. Oehlschlaeger

Department of Mechanical, Aerospace and Nuclear Engineering
Rensselaer Polytechnic Institute
Troy, New York

chowdm@rpi.edu and oehlsm@rpi.edu

MANE 4962 and 6962

Regular announcement

- 👉 HW 5 is due March 20, 2023
- 👉 Quiz 5 on March 20, 2023
- 👉 Office hour on Friday, March 17, 3-5 PM
- 👉 HW 5 discussion

→ Quiz 5 is based on Lectures 14 & 15

OR Gate and NOR Gate


input $\underline{x} = (x_1, x_2)$	output y	prediction
(0,0)	0	< 0.5
(0,1)	1	≥ 0.5
(1,0)	1	≥ 0.5
(1,1)	1	≥ 0.5

OR Gate: $w_0 = -1$ $w_1 = 2$ $w_2 = 2$

input $\underline{x} = (x_1, x_2)$	output y	prediction
(0,0)	1	≥ 0.5
(0,1)	0	< 0.5
(1,0)	0	< 0.5
(1,1)	0	< 0.5

NOR Gate: $w_0 = 1$ $w_1 = -2$ $w_2 = -2$

Note: w 's are not guaranteed to be unique



AND Gate and NAND Gate

input	output	prediction
$\underline{x} = (x_1, x_2)$	y	
(0,0)	0	< 0.5
(0,1)	0	< 0.5
(1,0)	0	< 0.5
(1,1)	1	≥ 0.5


AND Gate: $w_0 = -3$ $w_1 = 2$ $w_2 = 2$

input	output	prediction
$\underline{x} = (x_1, x_2)$	y	
(0,0)	1	≥ 0.5
(0,1)	0	< 0.5
(1,0)	0	< 0.5
(1,1)	0	< 0.5

NAND Gate: $w_0 = 3$ $w_1 = -2$ $w_2 = -2$

Note: w 's are not guaranteed to be unique

red-1, blue-0




XOR Gate

input $\underline{x} = (x_1, x_2)$	output y	prediction
(0,0)	0	×
(0,1)	1	×
(1,0)	1	×
(1,1)	0	×

XOR Gate: **No solution from logistic regression!**

- 👉 Use non-linear feature transformation. (The kernel trick) SVMs.
- 👉 Add extra layers. (Deep neural nets)



Support Vector Machines

A powerful and versatile machine learning model first proposed by Vladimir Vapnik. Check the very first notebook where we used SVM blindly to classify iris dataset which achieved very high classification accuracy even without data pre-processing.

- 👉 Linear or nonlinear classification
- 👉 Regression
- 👉 Outlier detection

On solving XOR Gate

```


1  import numpy as np
2  from sklearn import svm
3  X = np.array([[0, 0],
4               [0, 1],
5               [1, 0],
6               [1, 1]])
7  y = np.array([0, 1, 1, 0])
8  model = svm.SVC(kernel='rbf', C=1, gamma='auto')
9  model.fit(X, y)
10 predictions = model.predict(X)
11 print("Predictions: ", predictions) # should classify all the data points correctly

```

1958
Perceptron Rosenblatt
10 yr. Minsky &
Paper

1968-1980
MNIST digits
classified SVM


Linear Support Vector Machine Classifier



Adding more training data points within the margins of the classes will not change the classifier. The marginal data points are known as support vectors.

r

Support Vector Machine



Decision rule

$$\vec{w} \cdot \vec{x}_+ + b > 1, \quad y_i = +1$$

$$\vec{w} \cdot \vec{x}_- + b \leq -1, \quad y_i = -1$$

Width maximization

$$W_D = (\vec{x}_+ - \vec{x}_-) \cdot \frac{\vec{w}}{\|\vec{w}\|} = \frac{2}{\|\vec{w}\|}$$

$$\max\left(\frac{2}{\|\vec{w}\|}\right) \rightarrow \min(\|\vec{w}\|) \rightarrow \min\left(\frac{1}{2}\|\vec{w}\|^2\right)$$

Max^m margin classifier

Constrained optimization problem

$$L = \frac{1}{2}\|\vec{w}\|^2 - \sum_{i=1}^m \alpha_i [y_i(\vec{w} \cdot \vec{x}_i + b) - 1]$$

(Primal form)

$$L = \sum_{i=1}^m \alpha_i - \sum_{i=1}^m \sum_{j=1}^m \alpha_i \alpha_j y_i y_j \Phi(x_i, x_j)$$

(Dual form)

SVM Dual form
kernel function

solved under the constraints via the Lagrange multipliers

$$\alpha_i \geq 0 \quad \text{and} \quad \alpha_i \leq C$$

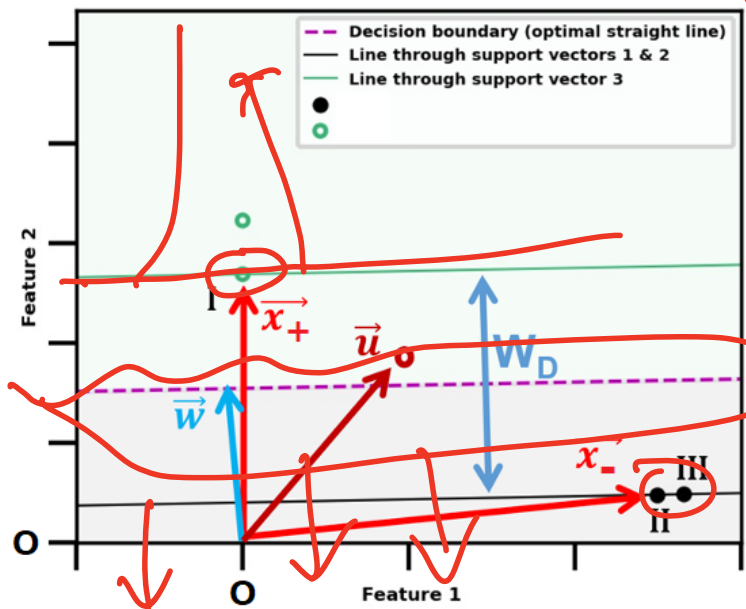
Larger values of C reduce misclassifications.

$$C = 1000$$

$$\vec{x}_+ + \vec{r} = \vec{x}_- \Rightarrow \vec{r} = \vec{x}_- - \vec{x}_+$$

$$\vec{r} \cdot \frac{\vec{w}}{\|\vec{w}\|} = W_D$$

Support Vector Machine



Decision rule

$$\vec{w} \cdot \vec{x}_+ + b \geq 1, \quad y_i = +1$$

$$\vec{w} \cdot \vec{x}_- + b \leq -1, \quad y_i = -1$$

Width maximization

$$W_D = (\vec{x}_+ - \vec{x}_-) \cdot \frac{\vec{w}}{\|\vec{w}\|} = \frac{2}{\|\vec{w}\|}$$

$$\max\left(\frac{2}{\|\vec{w}\|}\right) \rightarrow \min(\|\vec{w}\|) \rightarrow \min\left(\frac{1}{2}\|\vec{w}\|^2\right)$$

Contd.

Kernel functions ($\Phi(x_i, x_j)$)

$$\Phi(x_i, x_j) = x_i \cdot x_j \text{ (linear)}$$

$$\Phi(x_i, x_j) = e^{-\frac{\|x_i - x_j\|^2}{2\sigma^2}} = e^{-\gamma\|x_i - x_j\|^2} \text{ (radial basis function)}$$

$$\Phi(x_i, x_j) = (x_i \cdot x_j + k)^d \text{ (polynomial)}$$

σ relates sensitivity to variance in the feature vectors

$$\gamma = \frac{1}{2\sigma^2} \geq 0$$

Final solution

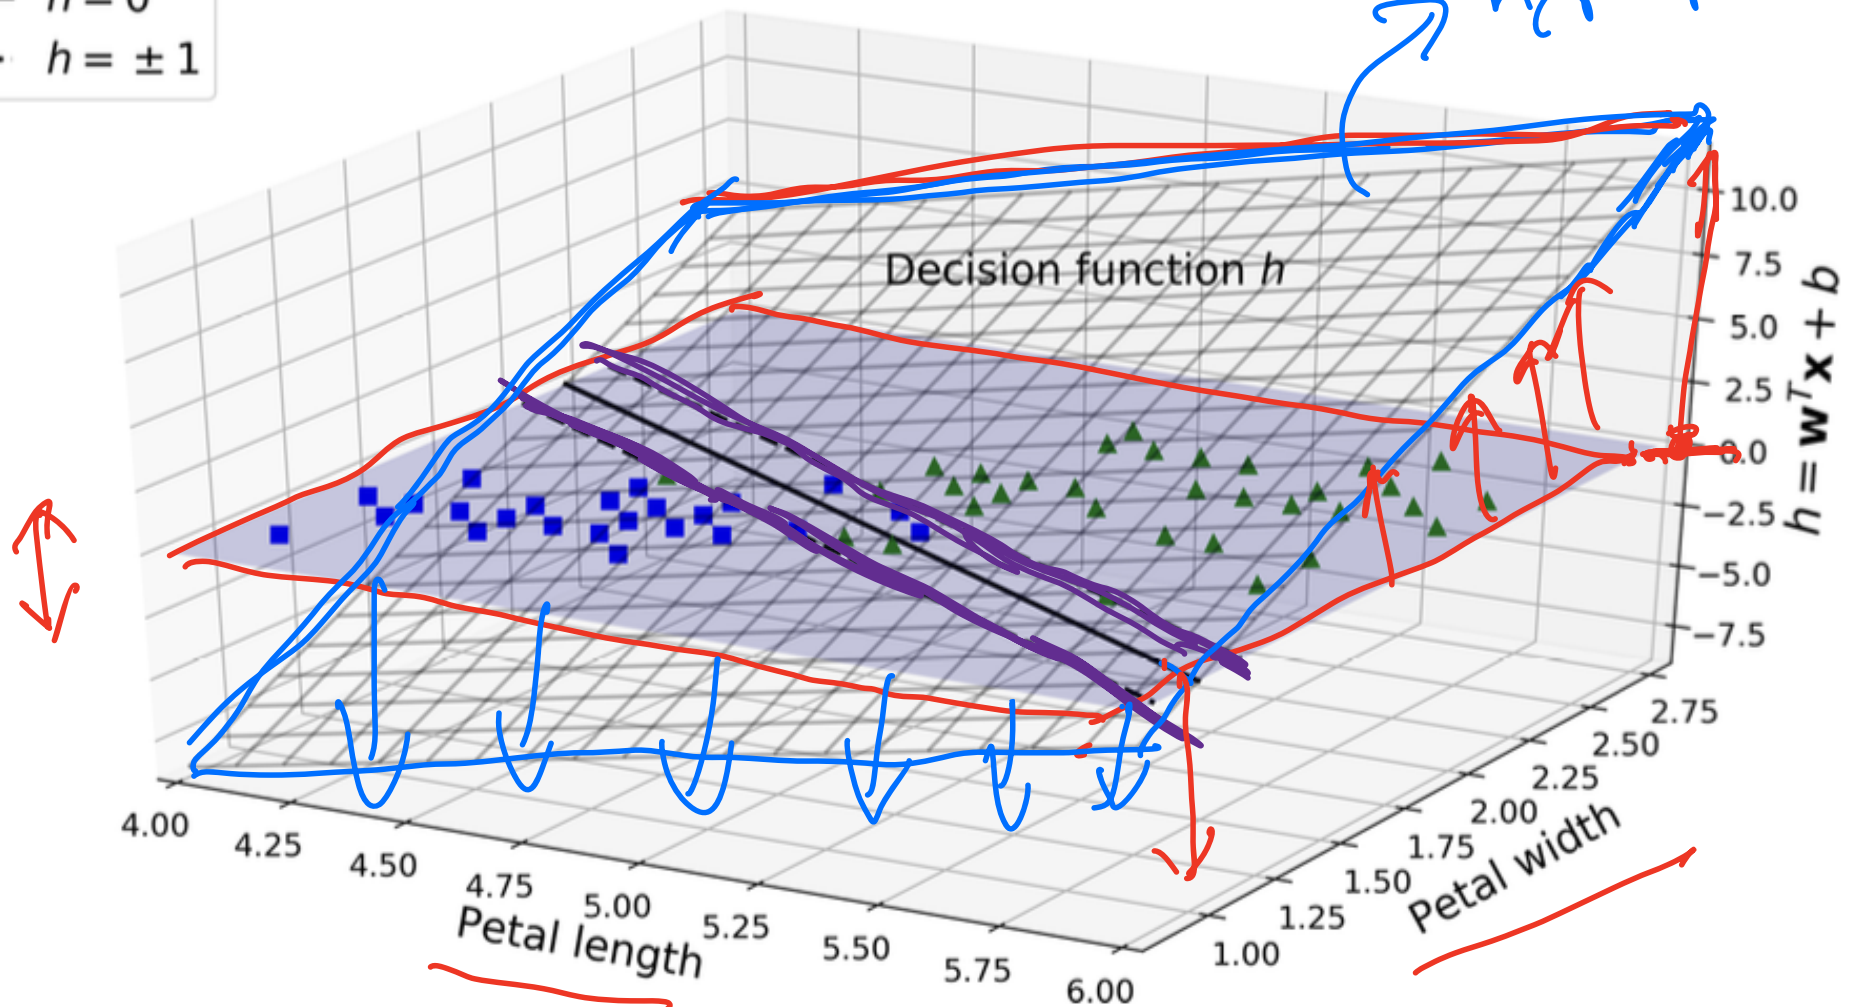
$$\vec{w} = \sum_{i=1}^m \alpha_i y_i \vec{u}$$

$$b = \frac{1}{N_S} \sum_{i \in y_i} y_i - \sum_{j \in y_j} \alpha_j y_j (x_i \cdot x_j)$$

$N_S \leq \text{Number of SV.}$


Under the hood

— $h = 0$
 --- $h = \pm 1$



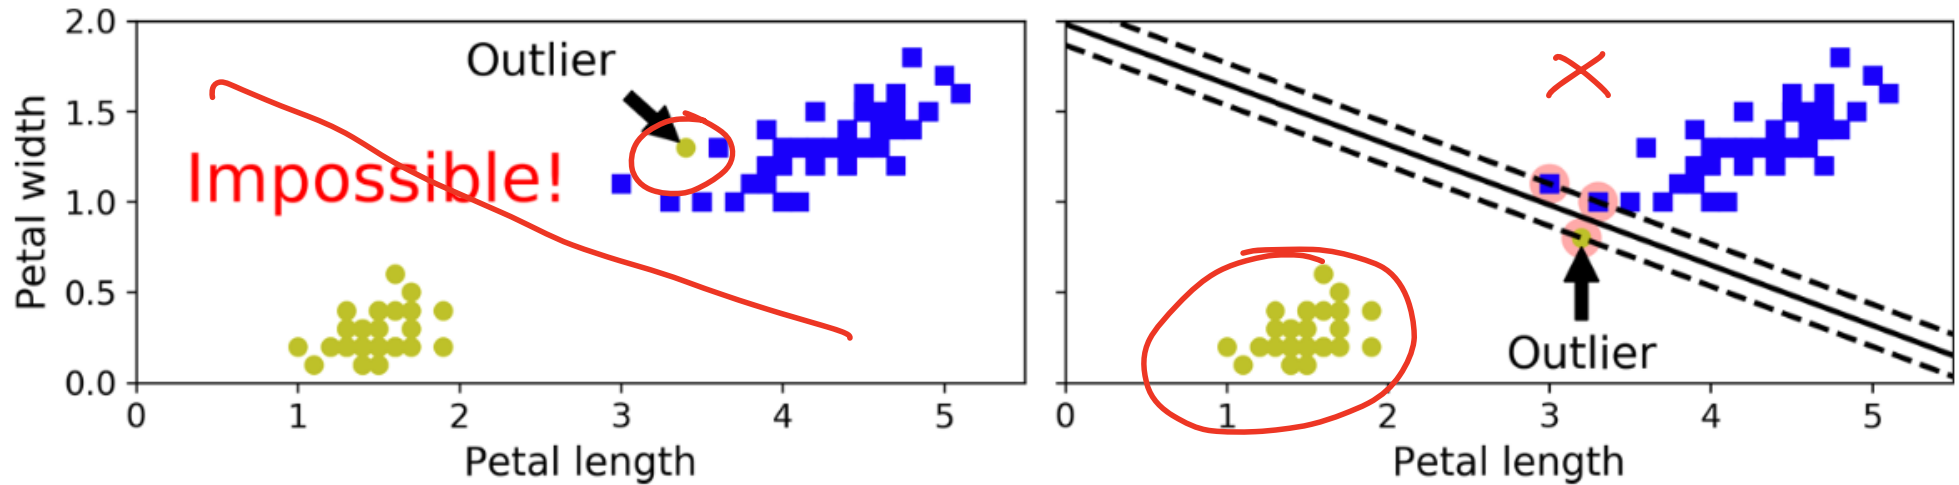
SVM + PCA (Dimensionality)

Feature scaling





Feature scaling will help widen the gap between the margins.

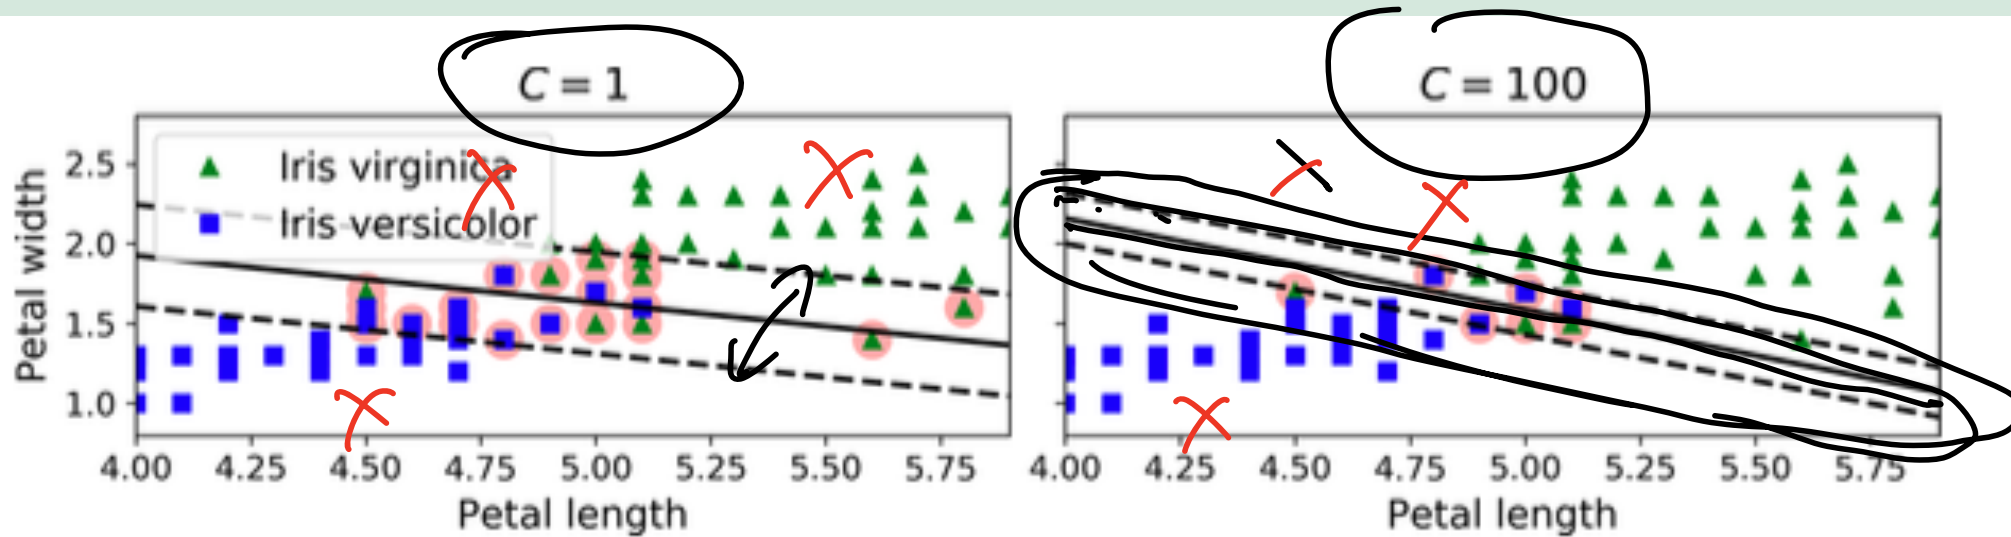
Hard margin vs Soft margin



Hard margin works only for linearly separable data. It strictly separates the two classes. Soft margin allows misclassification on either side of the margin.



Hard margin vs Soft margin




Margin violations are not desirable and thus minimized. Low C models generalize better. C is a regularization hyperparameter, called soft margin constant.

SVM

$C \rightarrow$ Reduce

high $C \rightarrow$ Overfit

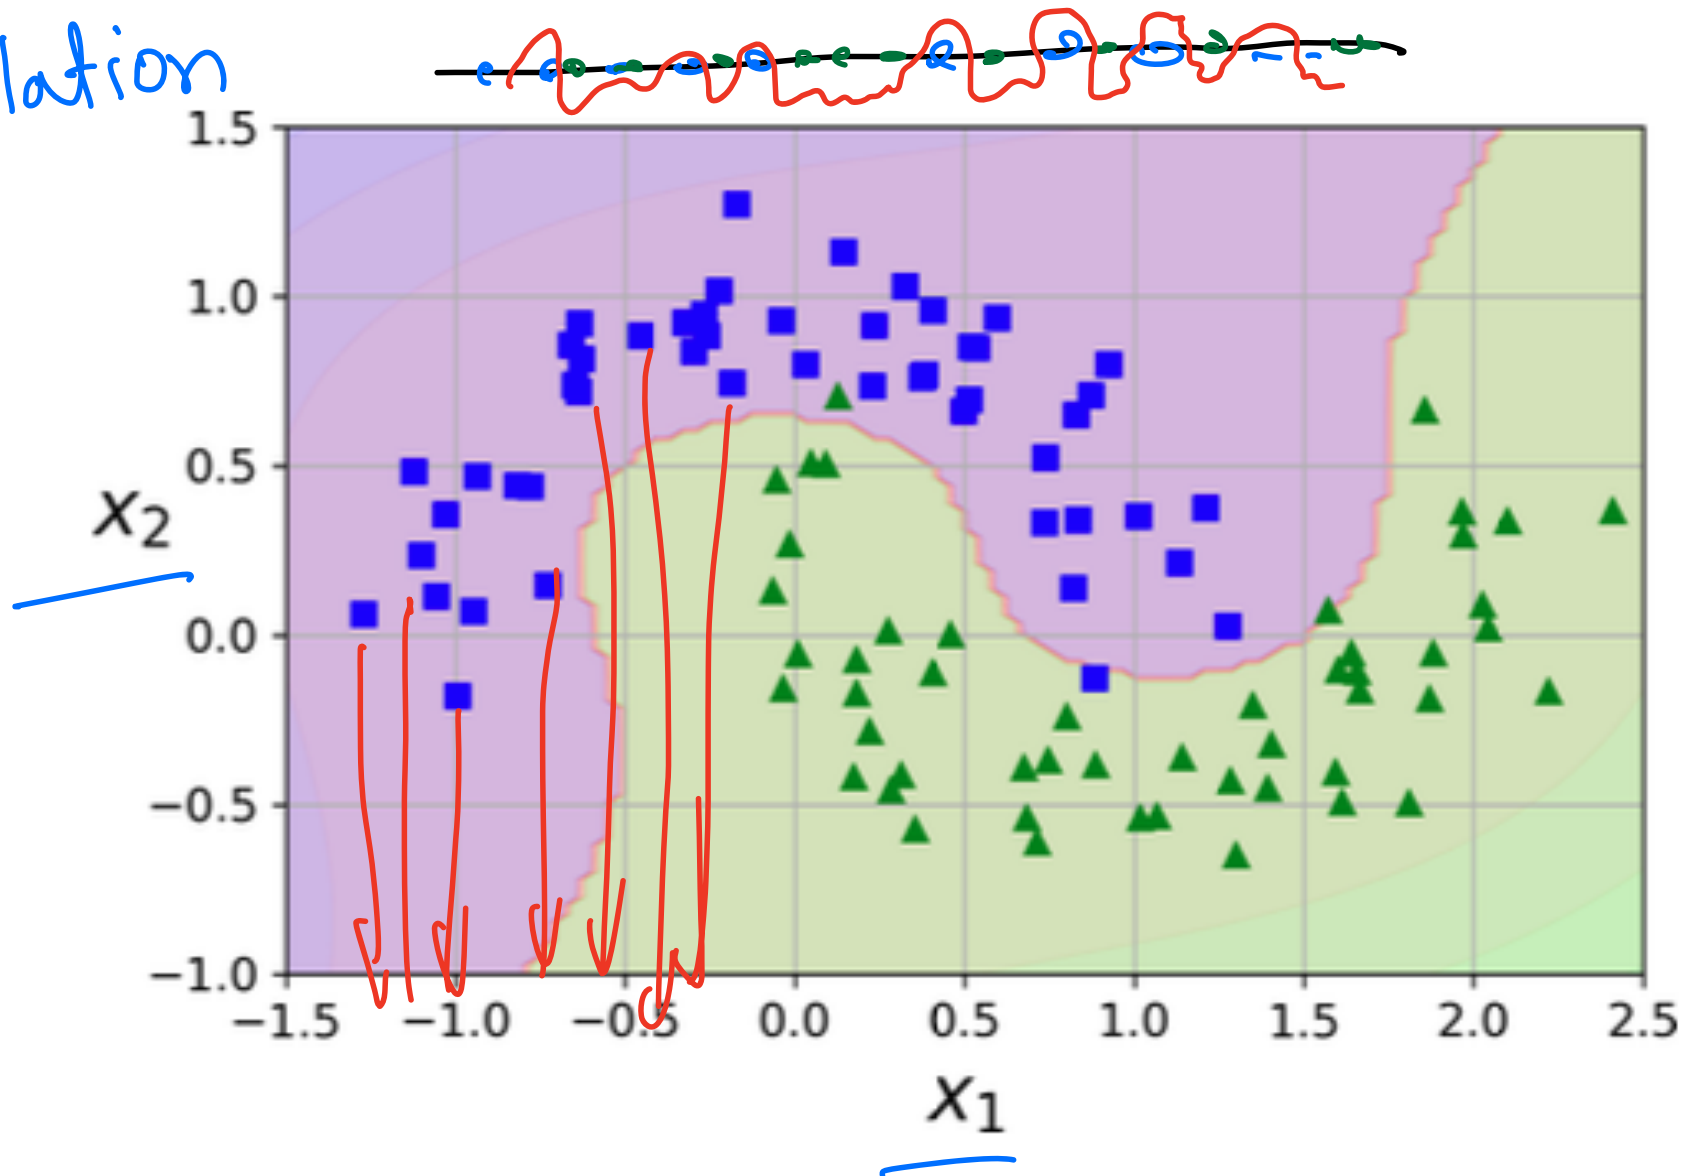
Non-linear SVM: Adding polynomial features



May make a problem linearly separable $x_2 = x_1^2$

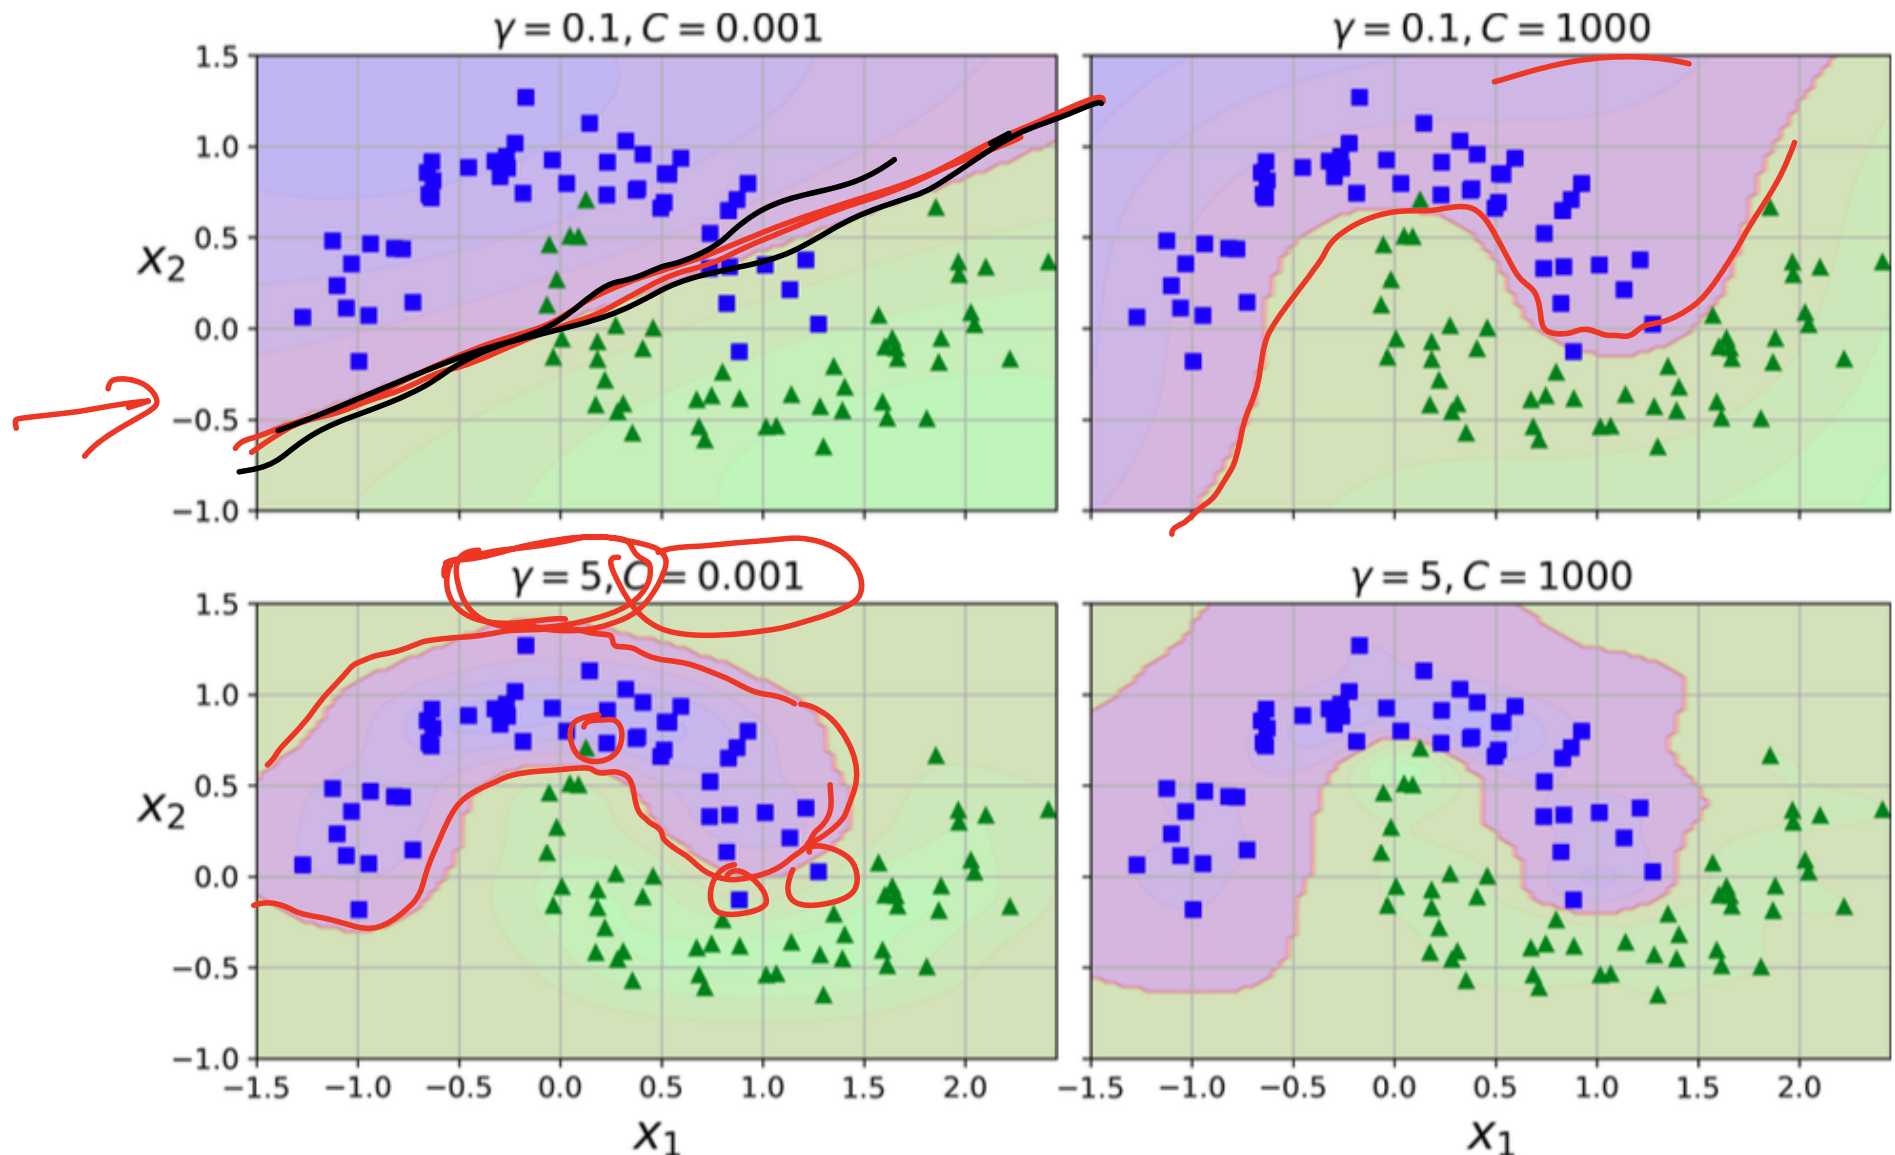
Kernel \rightarrow transform input features to high-D space

Ablation



Imagine classifying this dataset with just x_1 or x_2 alone.


Radial Basis Function Kernel $\rightarrow \phi(x_i, x_j)$



Heuristics for choosing the right kernel

Most commonly used kernels are linear, polynomial and radial basis function (Gaussian).

- Always try the linear kernel first, especially if the training set is very large or if it has plenty of features linear kernel will be faster.
- If the training set is not too large, you should also try the Gaussian RBF kernel; it works well in most cases.
- Experiment with other kernels, using cross-validation and grid search if additional computing resources and project time is available.



Scikit's implementation of SVM

Make sure to normalize the data. StandardScaler() is convenient for normalization. LinearSVC is faster but you must specify, the loss function to be hinge


```
1 from sklearn.pipeline import make_pipeline
2 from sklearn.preprocessing import StandardScaler
3 from sklearn.svm import SVC
4 clf = make_pipeline(StandardScaler(), SVC(gamma='auto'))
5 clf.fit(X, y)
```

Linear, RBF, polynomial

```
1 import numpy as np
2 from sklearn import datasets
3 from sklearn.pipeline import Pipeline
4 from sklearn.preprocessing import StandardScaler
5 from sklearn.svm import LinearSVC
6 iris = datasets.load_iris()
7
8 X = iris["data"][:, (2, 3)] # petal length, petal width
9 y = (iris["target"] == 2).astype(np.float64) # Iris virginica
10 svm_clf = Pipeline([
11     ("scaler", StandardScaler()),
12     ("linear_svc", LinearSVC(C=1, loss="hinge")),
13 ])
14 svm_clf.fit(X, y)
```

Linear

Hinge loss

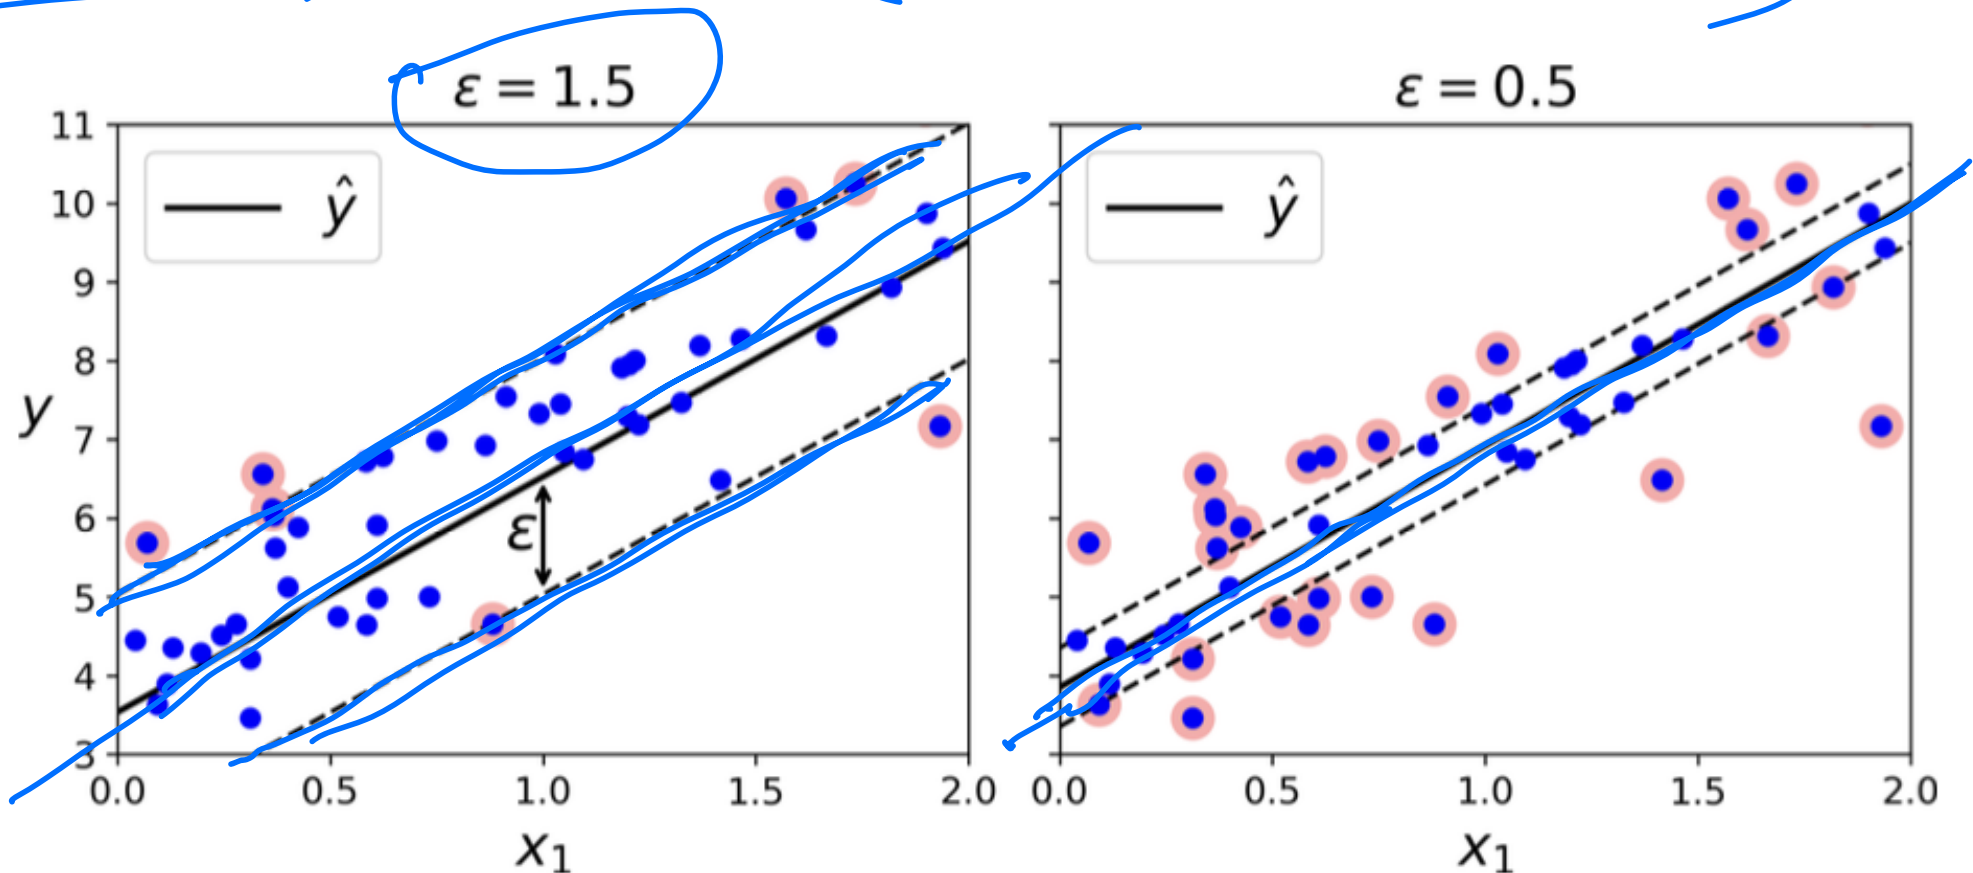


Computational Time

- Scikit's LinearSVC class is based on the LIBLINEAR library, which implements an optimized algorithm for linear SVMs.
- It does not support the kernel trick.
- Computation time scales almost linearly with the number of training data points and the number of features.
- The LIBLINEAR based linear SVM algorithm can be computationally expensive if you need very high precision. This is controlled by the tolerance hyperparameter ϵ (called tol in Scikit-Learn).
- SVC class is based on the LIBSVM library which supports the kernel trick.
- Computation time will scale as a product of either the square or the cubic order of the number of training data points and the number of features. So it will be excruciatingly slow for large datasets.
- Works well with sparse features (meaning most of the features have zero value). This makes the algorithm compute faster.

$$\overline{x_1, x_2, x_3, \dots, x_n}$$

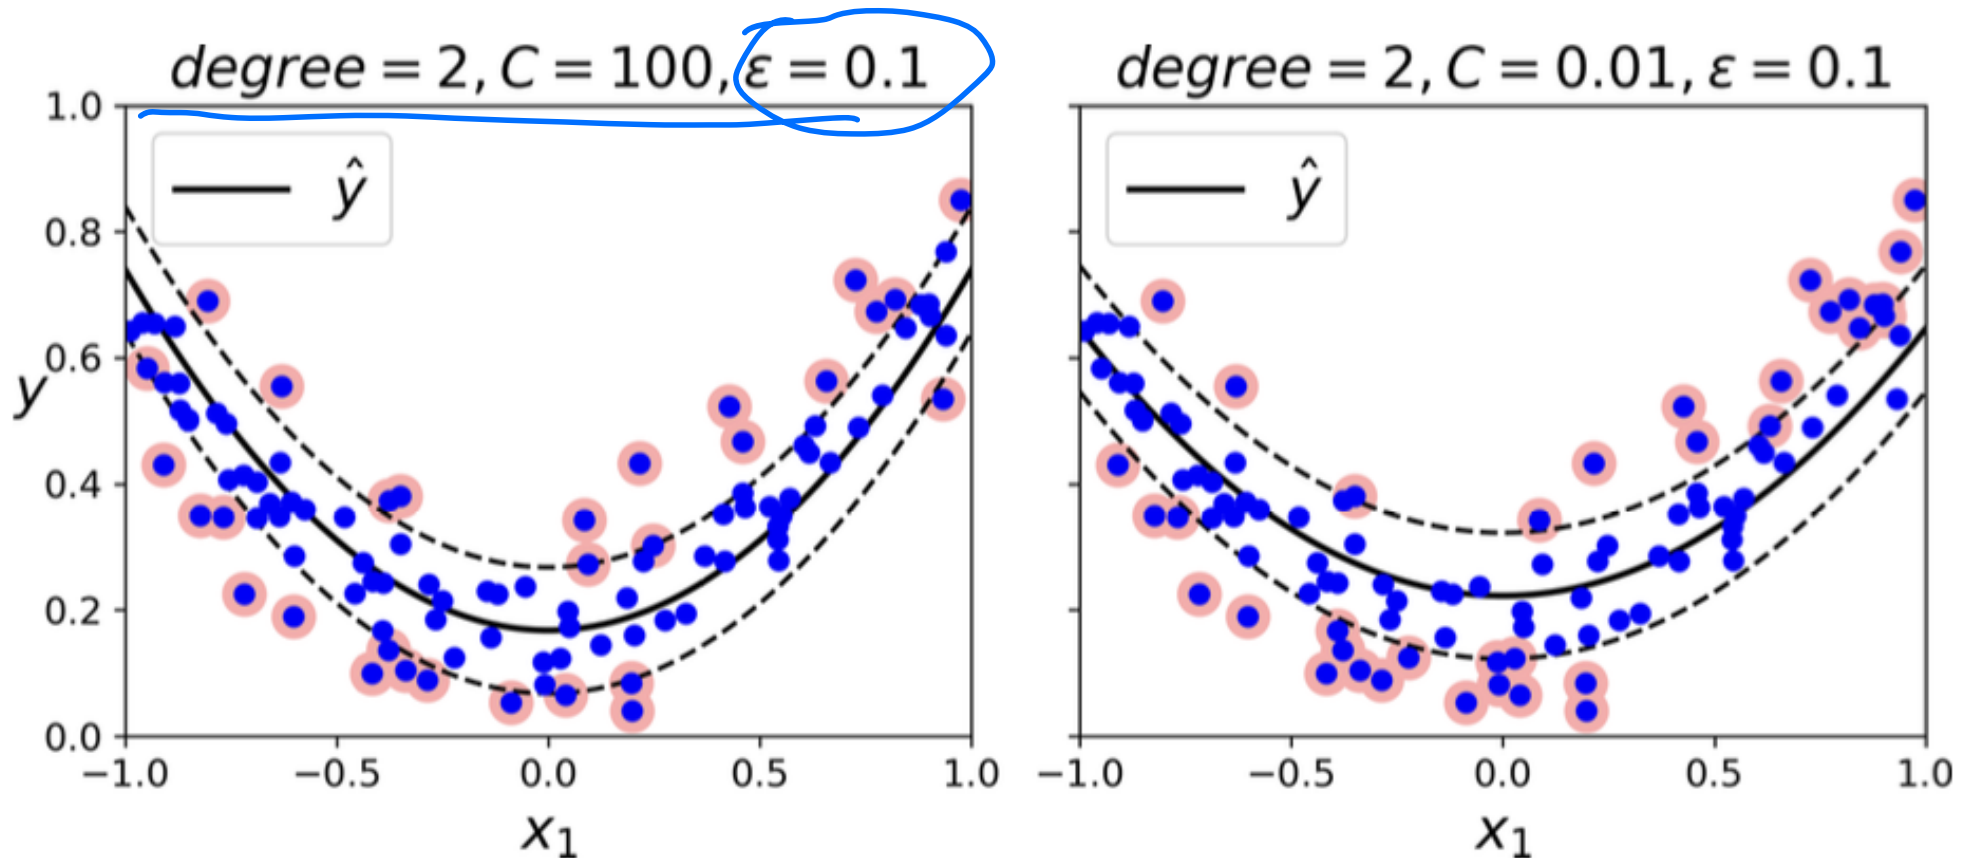
SVM Regression (Do not Use it)



SVM Regression attempts to fit as many data points as possible on the street while limiting margin violations. The width of the street is controlled by the epsilon hyperparameter. Results shown for linear kernel.

ϵ -insensitive

SVM Regression with polynomial kernel



```
1 from sklearn.svm import LinearSVR
2 svm_reg = LinearSVR(epsilon=1.5)
```

Advantages and disadvantages

Advantages:


- Works well for small to moderate datasets. It will work on larger dataset but computational time may be very high.
- Applicable to classification, regression, outlier detection.
- Needs only the support vectors to make decision.

Disadvantages:

- Can be slow and scaling issues may arise.
- Binary classifier by default, need to extend by training multiple classifiers
- For high-dimensional data (>3), interpretation is difficult.
- Unlike logistic regression, does not output estimated probabilities
- **Optimizing C is crucial for SVM performance.**


 Last Words = Dimensionality Reduction

Ini>



160,000

gase 10,000 →